Markov Chains

ISE/OR 560 Fall 2022

Hong Wan

September 19, 2022

North Carolina State University

NC STATE UNIVERSITY

System that evolves randomly in time..

A stochastic process is a collection of random variables $\{X(\tau), \tau \in T\}$, indexed by the parameter τ taking values in the parameter set T. The random variables take values in the set S, called the **state-space** of the stochastic process.

- Suppose we observe this system at discrete time points $n = 0, 1, 2, \cdots$. Let X_n be the state of the system at time n. We say that $\{X_n, n \geq 0\}$ is a *discrete-time stochastic process*

- If the system is observed continuously in time, with $X(t)$ being its state at time t, then $\{X(t), t \geq 0\}$ is described by a continuous time stochastic process.

parameter set

Sample Path

(a) Continuous-time, discrete state-space

3

(b) Continuous-time, continuous state-space

(c) Discrete-time, discrete state-space

- Queues: Let $X(t)$ be the number of customers waiting for service in a service facility such as an outpatient clinic.

- Supply Chains. Consider a supply chain of computer printers with three levels: the manufacturer (level 1), the regional warehouse (level 2), and the retail store (level 3). The printers are stored at all three levels. Let $X_i(t) =$ the number of printers at level i, $(1 \leq i \leq 3)$.

- A random walk is an example of a stochastic process in discrete time, where a particle starts at the origin at time 0 and moves one distance left with probability p or one distance unit right with probability $1p$ at each time unit.

- A single random variable X is completely described by its cumulative distribution function (cdf)

$$F(x) = \mathrm{P}(X \leq x), \quad -\infty < x < \infty.$$

- A multivariate random variable $(X_1, X_2, \cdots, X_n)$ is completely described by its joint cdf

$$F(x_1, x_2, \cdots, x_n) = \mathrm{P}(X_1 \leq x_1, X_2 \leq x_2, \cdots, X_n \leq x_n),$$

for all $-\infty < x_i < \infty$ and $i = 1, 2, \cdots, n$. If the parameter set T is finite, the stochastic process $\{X(\tau), \tau \in T\}$ is a multivariate random variable, and hence is completely described by the joint cdf.

- what about the case when T is not finite?

Suppose the sample paths of $\{X(t), t \geq 0\}$ are, with probability 1, right continuous with left limits, i.e.,

$$\lim_{s \downarrow t} X(s) = X(t), \quad \rbrace \; \text{stable behaviour}$$

and $\lim_{s \uparrow t} X(s)$ exists for each t. Furthermore, suppose the sample paths have a finite number of discontinuities in a finite interval of time with probability one. Then $\{X(t), t \geq 0\}$ is completely described $\rightarrow$ Stochastic process by a consistent family of finite dimensional joint cdfs

$$F_{t_1, t_2, \cdots, t_n}(x_1, x_2, \cdots, x_n) = P(X(t_1) \leq x_1, X(t_2) \leq x_2, \cdots, X(t_n) \leq x_n),$$

for all $-\infty < x_i < \infty, i = 1, \cdots, n, n \geq 1$ and all $0 \leq t_1 < t_2 < \cdots < t_n$.

joint distribution of multivariable random variable.

Let $\{X_n, n \geq 1\}$ be be a sequence of independent and identically distributed (iid) random variables with common distribution $F(\cdot)$. Define

[handwritten: don't move]

[handwritten: $S_1 = x_1$ I moved at step x_1]

$$S_0 = 0, S_n = X_1 + \cdots + X_n, \quad n \geq 1.$$

[handwritten: where I am at time n. Cumulative random walk.]

The stochastic process $\{S_n, n \geq 0\}$ is called a random walk. It is also completely characterized by $F(\cdot)$, since the joint distribution of $(S_0, S_1, \cdots, S_n)$ is completely determined by that of $(X_1, X_2, \cdots, X_n)$, which, is determined by $F(\cdot)$ that is defined as

$$F_n\left(X_1, X_2, \cdots, X_n\right) = \prod_{i=1}^{n} F\left(x_i\right),$$

[handwritten: because our random variables are independent.]

for all $-\infty < x_i < \infty, i = 1, 2, \cdots, n$, and $n \geq 1$.

Consider a system that is modeled by a discrete-time stochastic process $\{X_n, n \geq 0\}$ with a countable state-space S, say $\{0, 1, 2, \cdots\}$. Consider a fixed value of n that we shall call **"the present time"** or just the "present." Then X_n is called **the present (state) of the system**, $\{X_0, X_1, \cdots, X_{n-1}\}$ is called **the past of the system**, and $\{X_{n+1}, X_{n+2}, \cdots\}$ is called **the future of the system**. If $X_n = i$ and $X_{n+1} = j$, we say that the system has **jumped (or made a transition)** from state i to state j from time n to $n+1$.

- If the present state of the system is known, the future of the system is independent of its past. (is only dependent of where I am now, not yesterday)

- Intuition: the past affects the future only through the present.

$$\hookrightarrow P(X_{n+1} = j \mid X_n = i) = P(X_{n+1} = j \mid X_n, X_{n-1}, X_{n-2}, \ldots, X_0)$$

Definition of Discrete-Time Markov Chain: A stochastic process $\{X_n, n \geq 0\}$ with countable state-space S is called a DTMC if

System is Markovian Markov

1. for all $n \geq 0, X_n \in S$, $\rightarrow$ sample space do not changes

2. for all $n \geq 0$, and $i, j \in S$, then $P(X_{n+1} = j \mid X_n = i,$
$$X_{n-1}, X_{n-2}, \ldots, X_0) = P(X_{n+1} = j \mid X_n = i)$$

what is the meaning of this in words?

$\rightarrow$ It do not changes over time.

Time-Homogeneous DTMC. A DTMC $\{X_n, n \geq 0\}$ with countable state-space S is said to be time-homogeneous if

$$P(x_{n+1} = j \mid x_n = i) = P_{ij} \quad \forall \quad n \geqslant 0, \; i, j \in S.$$

Probability of going from state i to state j.

$\hookrightarrow$ is the probability always the same?

Let

$$P = \left[p_{i,j} \right]$$

denote the matrix of the conditional probabilities $p_{i,j}$. We call $p_{i,j}$ the transition probability from state i to state j. The matrix P is called the one-step transition probability matrix or just the transition probability matrix. When S is finite, say $S = \{1, 2, \cdots, m\}$, one can display P as a matrix as follows:

$$P = \begin{bmatrix} p_{1,1} & p_{1,2} & \cdots & p_{1,m-1} & p_{1,m} \\ p_{2,1} & p_{2,2} & \cdots & p_{2,m-1} & p_{2,m} \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ p_{m-1,1} & p_{m-1,2} & \cdots & p_{m-1,m-1} & p_{m-1,m} \\ p_{m,1} & p_{m,2} & \cdots & p_{m,m-1} & p_{m,m} \end{bmatrix}$$

$\}$ Transition matrix

- needs to be square
- all $p_{ij} \geq 0$
- $p_{ij} \neq p_{ji}$ all the time

$|$ can be used as transition matrix.

13

Stochastic Matrix. A square matrix $P = \left[p_{i,j} \right]$ is called stochastic if

(i). $p_{i,j} \geq 0$ for all $i, j \in S$,

(ii). $\displaystyle\sum_{j \in S} p_{i,j} = 1$ for all $i \in S$. rows

Stochastic Matrix. A square matrix $P = \left[p_{i,j}\right]$ is called stochastic if

$$\text{(i).} \quad p_{i,j} \geq 0 \quad \text{for all } i, j \in S,$$

$$\text{(ii).} \quad \sum_{j \in S} p_{i,j} = 1 \quad \text{for all } i \in S.$$

how do I interpret.

We can prove that the one-step transition probability matrix of a DTMC is stochastic.

Suppose we specify the distribution of X_0 externally. Let

$$a_i = \mathrm{P}\left(X_0 = i\right), \quad i \in S,$$

and

$$a = [a_i]_{i \in S}$$

be a row vector representing the probability mass function (pmf) of X_0. We say that a is the initial distribution of the DTMC.

Theorem: A DTMC $\{X_n, n \geq 0\}$ is completely described by its initial distribution a and the transition probability matrix P.

(x_0)

Transition Diagram. Consider a DTMC $\{X_n, n \geq 0\}$ on state-space $\{1, 2, 3\}$ with the following transition probability matrix:

$$P = \begin{bmatrix} .1 & .2 & .7 \\ .6 & 0 & .4 \\ .4 & 0 & .6 \end{bmatrix}$$

Probability Calculation

Joint Distributions. Let $\{X_n, n \geq 0\}$ be a DTMC on state-space $\{1, 2, 3, 4\}$ with the transition probability matrix given below:

$$P = \begin{bmatrix} 0.1 & 0.2 & 0.3 & 0.4 \\ 0.2 & 0.2 & 0.3 & 0.3 \\ 0.5 & 0.0 & 0.5 & 0.0 \\ 0.6 & 0.2 & 0.1 & 0.1 \end{bmatrix}$$

The initial distribution is

$$P(x_n = 0.25) \mid n = 1, 2, 3, 4$$

$$a = \begin{bmatrix} 0.25 & 0.25 & 0.25 & 0.25 \end{bmatrix}. \Rightarrow \text{probability of being in a state.}$$

1. Compute $P(X_3 = 4, X_2 = 1, X_1 = 3, X_0 = 1) = (0.25) \cdot P_{1,13} \cdot P_{3,1} \cdot P_{1,4}$

2. Compute $P(X_3 = 4, X_2 = 1, X_1 = 3)$.

· the longer the sequence, the smaller the probability.

Two-State DTMC. One of the simplest DTMCs is one with two states, labeled 1 and 2. Thus $S = \{1, 2\}$. Such a DTMC has a transition matrix as follows:

$$\begin{bmatrix} \alpha & 1 - \alpha \\ 1 - \beta & \beta \end{bmatrix},$$

where $0 \leq \alpha, \beta \leq 1$.

The transition diagram:

Consider a simple weather model in which we classify the day's weather as either "sunny" or "rainy." On the basis of previous data we have determined that if it is sunny today, there is an 80% chance that it will be sunny tomorrow regardless of the past weather; whereas, if it is rainy today, there is a 30% chance that it will be rainy tomorrow, regardless of the past.

Weather Model

Suppose two drugs are available to treat a particular disease, and we need to determine which of the two drugs is more effective. This is generally accomplished by conducting clinical trials of the two drugs on actual patients. Here we describe a clinical trial setup that is useful if the response of a patient to the administered drug is sufficiently quick, and can be classified as "effective" or "ineffective." Suppose drug i is effective with probability $p_i, i = 1, 2$. In practice the values of p_1 and p_2 are unknown, and the aim is to determine if $p_1 \geq p_2$ or $p_2 \geq p_1$. Ethical reasons compel us to use the better drug on more patients. This is achieved by using the play the winner rule as follows.

Clinical Trials

The initial patient (indexed as patient 0) is given either drug 1 or 2 at random. If the nth patient is given drug $i(i = 1, 2)$ and it is observed to be effective for that patient, then the same drug is given to the $(n + 1)$-st patient; if it is observed to be ineffective then the $(n + 1)$-st patient is given the other drug. Thus we stick with a drug as long as its results are good; when we get a bad result, we switch to the other drug - hence the name "play the winner."

Consider two gamblers, A and B, who have a combined fortune of N dollars. They bet one dollar each on the toss of a coin. If the coin turns up heads, *A* wins a dollar from *B*, and if the coin turns up tails, *B* wins a dollar from A. Suppose the successive coin tosses are independent, and the coin turns up heads with probability p and tails with probability $q = 1 - p$. The game ends when either A or B is broke (or ruined).

random walk.

absorbing state.

28

Gambler's Ruin

Consider two urns labeled A and B, containing a total of N white balls and N red balls among them. An experiment consists of picking one ball at random from each urn and interchanging them. This experiment is repeated in an independent fashion. Let X_n be the number of white balls in urn A after n repetitions of the experiment. Assume that initially urn A contains all the white balls, and urn B contains all the red balls. Thus $X_0 = N$. Note that X_n tells us precisely the contents of the two urns after n experiments: if $X_n = i$, urn A contains i white balls and $N - i$ red balls; and urn B contains $N - i$ white balls and i red balls.